In this lab I've got a local LLM running in Google Colab via Ollama and am connecting to it from Python using an OpenAI compatible client. I'm using Pydantic and the Instructor library to get the model to return properly structured JSON "malware cards". I've built a weekly cybersecurity news digest by pulling in the RSS feeds and having the model rank and summarize the top stories with an impact score.

The key thing I've learned is that you can't just take AI output on its own trustworthiness. The malware cards in particular had some rather incorrect or inconsistent information so they need to be checked against a more reliable source. I've also learned that Colab is good for practice but not really for actual SOC work given the time limits and the data privacy issues.

In a SOC this could be used to auto-generate your weekly threat intel briefs or to add a bit of context to your alert tickets. My next step would be to run a small pilot of the news digest on a local server with an analyst reviewing every single summary before it got out.