The big risk isn't coming up with the answer from nothing really, it's overconfidence. [X%] of the answers were fully sourced, [Y%] had at least one unsupported assertion (mostly on the [weak-retrieval / multi-hop / numeric] questions). The model is very good at making up the bits that aren't there.

To fix this we'd improve both retrieval and generation. For retrieval we'd do re-ranking and set a relevance cutoff so the worst passages don't even get to the model. For generation we'd require citations for every claim and give the model a clear 'I don't have enough information' option. We're also going to test on harder and out of scope questions and look at faithfulness as well as how often the model actually says it doesn't have enough information rather than just accuracy.

Evaluation has changed the way I think about readiness. The system appeared to be ready as its answers were good-sounding, but good-sounding is not the same as being reliable. I now have a clear benchmark (e.g. [≥95% faithfulness with correct refusals on unanswerable questions]), will only deploy once that benchmark is met and will continue to monitor post-launch with human review for high-stakes answers.