1. **Explain the lexical failure.** A misattribution keeps the passage exactly as written and simply swaps the entity, so ROUGE and BLEU will reward it.
2. **Cite your own evidence.** A is a clean example of good retrieval and bad reading, which directly supports scoring the two separately.
3. **Require span-level citations.** Someone could cite the correct passage and still be wrong. The cited span must support the specific entity.
4. **Gate on the worst case.** Require zero ungrounded answers on the critical items (not just a mean threshold). If you allow abstention, you need to track over-abstention as well.
5. **Calibrate the judge on entity swaps.** They're the toughest case for an LLM judge, so they should be in the human-labeled set.
6. **Reframe the harness item** as "the eval pipeline itself must be trustworthy".